== Metric definitions

gold_agreement is gold_hits divided by gold_items. a_side_bias is the share of left among an annotator's left and right verdicts on balanced comparisons, minus 0.5; balanced_verdicts counts those verdicts.

Agreement is computed over rated units: comparisons with at least two admissible verdicts. Both_bad folds into tie for agreement, and only for agreement. With n ratings in a unit and n_c in category c, observed_agreement is the sum over units of the sum over c of n_c(n_c-1), divided by the sum over units of n(n-1). rater_pairs is the sum of n(n-1)/2 and rated_units the number of units. The expected agreement is the sum over categories of the squared share of that category among all ratings in the units, and kappa is observed minus expected, divided by one minus expected (0 when expected is 1).

by_domain groups units by the prompt's domain. by_vendor forms a unit from one vendor's admissible verdicts on a rated unit when there are at least two of them.

Every reported ratio is quantised to six places, half to even, at the reporting step and at no earlier step. A ratio whose denominator is zero does not exist and is reported as null.